

Утверждаю:

Гапанюк Ю.Е.

"__" _____ 2019 г.

**Курсовая работа по курсу
“Технологии машинного обучения”
“Бинарная классификация”**

Вариант №_1_

Пояснительная записка
(вид документа)

писчая бумага
(вид носителя)

22
(количество листов)

ИСПОЛНИТЕЛЬ:

студент группы ИУ5-62
Андреев А.С.

"__" _____ 2019 г.

Содержание

1. Задание	3
2. Введение	4
3. Основная часть	5
a. Постановка задачи	5
b. Описание выбранного датасета (листинг)	5
c. Решение задачи регрессии (листинг)	12
d. Графическая реализация (листинг)	16
4. Заключение	21
5. Список литературы.....	22

Задание

В данной курсовой работе необходимо предпринять следующие шаги:

1. Поиск и выбор набора данных для построения моделей машинного обучения. На основе выбранного набора данных студент должен построить модели машинного обучения для решения или задачи классификации, или задачи регрессии.
2. Проведение разведочного анализа данных. Построение графиков, необходимых для понимания структуры данных. Анализ и заполнение пропусков в данных.
3. Выбор признаков, подходящих для построения моделей. Кодирование категориальных признаков Масштабирование данных. Формирование вспомогательных признаков, улучшающих качество моделей.
4. Проведение корреляционного анализа данных. Формирование промежуточных выводов о возможности построения моделей машинного обучения.

В зависимости от набора данных, порядок выполнения пунктов 2, 3, 4 может быть изменен.

5. Выбор метрик для последующей оценки качества моделей. Необходимо выбрать не менее трех метрик и обосновать выбор.
6. Выбор наиболее подходящих моделей для решения задачи классификации или регрессии. Необходимо использовать не менее пяти моделей, две из которых должны быть ансамблевыми.
7. Формирование обучающей и тестовой выборок на основе исходного набора данных.
8. Построение базового решения (baseline) для выбранных моделей без подбора гиперпараметров. Производится обучение моделей на основе обучающей выборки и оценка качества моделей на основе тестовой выборки.
9. Подбор гиперпараметров для выбранных моделей. Рекомендуется использовать методы кросс-валидации. В зависимости от используемой библиотеки можно применять функцию GridSearchCV, использовать перебор параметров в цикле, или использовать другие методы.
10. Повторение пункта 8 для найденных оптимальных значений гиперпараметров. Сравнение качества полученных моделей с качеством baseline-моделей.
11. Формирование выводов о качестве построенных моделей на основе выбранных метрик. Результаты сравнения качества рекомендуется отобразить в виде графиков и сделать выводы в форме текстового описания. Рекомендуется построение графиков обучения и валидации, влияния значений гиперпараметров на качество моделей и т.д.

Приведенная схема исследования является рекомендуемой. Возможно выполнение курсовой работы на нестандартную тему, которая должна быть предварительно согласована с ответственным за прием курсовой работы.

Введение

Курсовая работа – самостоятельная часть учебной дисциплины «Технологии машинного обучения» – учебная и практическая исследовательская студенческая работа, направленная на решение комплексной задачи машинного обучения. Результатом курсовой работы является отчет, содержащий описания моделей, тексты программ и результаты экспериментов.

Курсовая работа опирается на знания, умения и владения, полученные студентом в рамках лекций и лабораторных работ по дисциплине.

В рамках данной курсовой работы необходимо применить навыки, полученные в течение курса «Технологии машинного обучения», и обосновать полученные результаты.

Основная часть

Постановка задачи

Датасет представляет собой описание дома и его стоимость. В нем 19 признаков, среди которых количество спален, ванн, площадь, по которым предсказывается стоимость (это целевой признак). Это классическая задача регрессии.

Разделительная плоскость, описанная выше, была получена с использованием Multisurface Method-Tree (MSM-T) [К. П. Беннетт, "Построение дерева решений с помощью линейного программирования". Труды 4-го Среднего Запада Общества искусственного интеллекта и когнитивной науки, стр. 97-101, 1992], метод классификации, который использует линейное программирование для построения дерева решений. Соответствующие элементы были выбраны с использованием исчерпывающего поиска в пространстве 1-4 элементов и 1-3 разделительных плоскостей. Фактическая линейная программа, используемая для получения разделяющей плоскости в трехмерном пространстве, описана в: [К. П. Беннетт и О. Л. Мангасарян: «Надежное распознавание линейного программирования двух линейно неразделимых множеств», Методы оптимизации и программное обеспечение 1, 1992, 23-34].

Описание выбранного датасета

Информация об атрибутах:

id a notation for a house
date Date house was sold
price Price is prediction target
bedrooms Number of Bedrooms/House
bathrooms Number of bathrooms/House
sqft_living square footage of the home
sqft_lot square footage of the lot
floors Total floors (levels) in house
water frontHouse which has a view to a waterfront
viewHas been viewed
conditionHow good the condition is (Overall)
gradeoverall grade given to the housing unit, based on King County grading system
sqft_abovesquare footage of house apart from basement
sqft_basementsquare footage of the basement
yr_builtBuilt Year
yr_renovatedYear when house was renovated
zipcodezip
latLatitude coordinate
longLongitude coordinate
sqft_living15Living room area in 2015(implies-- some renovations) This might or might not have affected the lotsize area
sqft_lot15lotSize area in 2015(implies-- some renovations)

```
In [311]: import numpy as np
import pandas as pd
pd.set_option('display.max.rows', 1000)
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style='ticks')
```

1) Поиск и выбор набора данных для построения моделей машинного обучения.

```
In [312]: data = pd.read_csv('data/kc_house_data.csv')
```

House Sales in King County, USA

Predict house price using regression

<https://www.kaggle.com/harlfoxem/housesalesprediction>
(<https://www.kaggle.com/harlfoxem/housesalesprediction>)

This dataset contains house sale prices for King County, which includes Seattle. It includes homes sold between May 2014 and May 2015.

It's a great dataset for evaluating simple regression models.

19 house features plus the price and the id columns, along with 21613 observations:

- id – a notation for a house
- date – Date house was sold
- price – Price is prediction target
- bedrooms – Number of Bedrooms/House
- bathrooms – Number of bathrooms/House
- sqft_living – square footage of the home
- sqft_lot – square footage of the lot
- floors – Total floors (levels) in house
- waterfront – House which has a view to a waterfront
- view – Has been viewed
- condition – How good the condition is (Overall)
- grade – overall grade given to the housing unit, based on King County grading system
- sqft_above – square footage of house apart from basement
- sqft_basement – square footage of the basement
- yr_built – Built Year
- yr_renovated – Year when house was renovated
- zipcode – zip
- lat – Latitude coordinate
- long – Longitude coordinate
- sqft_living15 – Living room area in 2015(implies-- some renovations) This might or might not have affected the lotsize area
- sqft_lot15 – lotSize area in 2015(implies-- some renovations)

```
In [313]: data.head()
```

```
Out[313]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	

5 rows × 21 columns

```
In [314]: data.shape
```

```
Out[314]: (21613, 21)
```

```
In [315]: data.dtypes
```

```
Out[315]: id                int64
date                object
price              float64
bedrooms           int64
bathrooms          float64
sqft_living        int64
sqft_lot           int64
floors             float64
waterfront         int64
view               int64
condition          int64
grade              int64
sqft_above         int64
sqft_basement      int64
yr_built           int64
yr_renovated       int64
zipcode            int64
lat                float64
long               float64
sqft_living15      int64
sqft_lot15         int64
dtype: object
```

2) Проведение разведочного анализа данных. Построение графиков, необходимых для понимания структуры данных. Анализ и заполнение пропусков в данных.



```
In [316]: data.describe()
```

```
Out[316]:
```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lk
count	2.161300e+04	2.161300e+04	21613.000000	21613.000000	21613.000000	2.161300e+C
mean	4.580302e+09	5.400881e+05	3.370842	2.114757	2079.899736	1.510697e+C
std	2.876566e+09	3.671272e+05	0.930062	0.770163	918.440897	4.142051e+C
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+C
25%	2.123049e+09	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+C
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+C
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+C
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+C

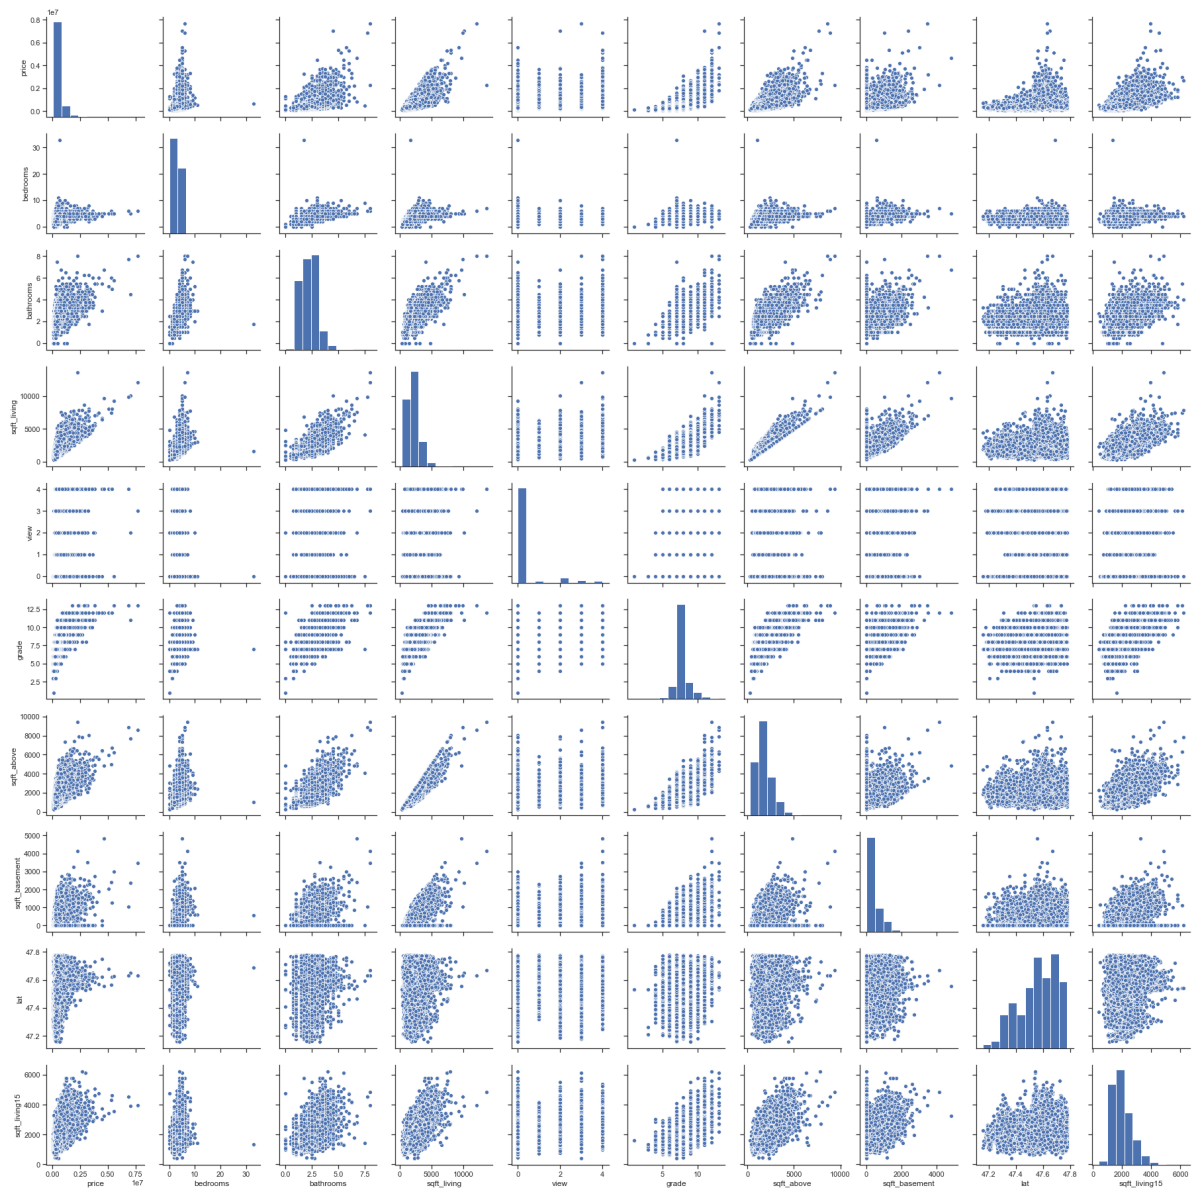
```
In [317]: fig, ax = plt.subplots(figsize=(10,10))
sns.distplot(data['price'])
```

```
Out[317]: <matplotlib.axes._subplots.AxesSubplot at 0x12d6d0748>
```



```
In [424]: sns.pairplot(data) # парные диаграммы
```

```
Out[424]: <seaborn.axisgrid.PairGrid at 0x1401c5278>
```



```
In [319]: data.isnull().sum()
```

```
Out[319]: id                0
          date              0
          price             0
          bedrooms          0
          bathrooms         0
          sqft_living       0
          sqft_lot          0
          floors            0
          waterfront        0
          view              0
          condition         0
          grade             0
          sqft_above        0
          sqft_basement     0
          yr_built          0
          yr_renovated      0
          zipcode           0
          lat               0
          long              0
          sqft_living15     0
          sqft_lot15        0
          dtype: int64
```

Пропусков нет, заполнять нечего.

3) Выбор признаков, подходящих для построения моделей. Кодирование категориальных признаков Масштабирование данных. Формирование вспомогательных признаков, улучшающих качество моделей.


In [320]: data.corr()

Out[320]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floor
id	1.000000	-0.016762	0.001286	0.005160	-0.012258	-0.132109	0.01852
price	-0.016762	1.000000	0.308350	0.525138	0.702035	0.089661	0.25679
bedrooms	0.001286	0.308350	1.000000	0.515884	0.576671	0.031703	0.17542
bathrooms	0.005160	0.525138	0.515884	1.000000	0.754665	0.087740	0.50065
sqft_living	-0.012258	0.702035	0.576671	0.754665	1.000000	0.172826	0.35394
sqft_lot	-0.132109	0.089661	0.031703	0.087740	0.172826	1.000000	-0.00520
floors	0.018525	0.256794	0.175429	0.500653	0.353949	-0.005201	1.00000
waterfront	-0.002721	0.266369	-0.006582	0.063744	0.103818	0.021604	0.02369
view	0.011592	0.397293	0.079532	0.187737	0.284611	0.074710	0.02944
condition	-0.023783	0.036362	0.028472	-0.124982	-0.058753	-0.008958	-0.26376
grade	0.008130	0.667434	0.356967	0.664983	0.762704	0.113621	0.45818
sqft_above	-0.010842	0.605567	0.477600	0.685342	0.876597	0.183512	0.52388
sqft_basement	-0.005151	0.323816	0.303093	0.283770	0.435043	0.015286	-0.24570
yr_built	0.021380	0.054012	0.154178	0.506019	0.318049	0.053080	0.48931
yr_renovated	-0.016907	0.126434	0.018841	0.050739	0.055363	0.007644	0.00633
zipcode	-0.008224	-0.053203	-0.152668	-0.203866	-0.199430	-0.129574	-0.05912
lat	-0.001891	0.307003	-0.008931	0.024573	0.052529	-0.085683	0.04961
long	0.020799	0.021626	0.129473	0.223042	0.240223	0.229521	0.12541
sqft_living15	-0.002901	0.585379	0.391638	0.568634	0.756420	0.144608	0.27988
sqft_lot15	-0.138798	0.082447	0.029244	0.087175	0.183286	0.718557	-0.01126


```
In [323]: plt.figure(figsize=(16, 16))
sns.heatmap(data.corr(), annot=True, fmt='.3f')
```

```
Out[323]: <matplotlib.axes._subplots.AxesSubplot at 0x143548908>
```



```
In [324]: # дробнем коллинеарные фичи:
# data = data.drop(columns=['grade', 'sqft_above', 'bathrooms'])
```

```
In [325]: plt.figure(figsize=(16, 16))  
sns.heatmap(data.corr(), annot=True, fmt='.3f')
```

```
Out[325]: <matplotlib.axes._subplots.AxesSubplot at 0x142f6f748>
```



Кодирование категориальных признаков

```
In [326]: cat_cols = []  
for col in data.columns:  
    if data[col].dtype == 'object':  
        cat_cols.append(col)  
cat_cols
```

```
Out[326]: ['date']
```

```
In [327]: for col in cat_cols:
           print('{}: {} unique values'.format(col, len(data[col].unique(
))))
```

```
`date`: 372 unique values
```

```
In [328]: from sklearn.preprocessing import LabelEncoder
```

```
In [329]: encoding_of_cat = {}
           for col in cat_cols:
               le = LabelEncoder()
               data[[col]] = le.fit_transform(data[col])
               encoding_of_cat[col] = le
```

```
In [330]: data.corr()
```

Out[330]:

	date	price	bedrooms	bathrooms	sqft_living	view	grad
date	1.000000	-0.004649	-0.016964	-0.034481	-0.034570	-0.001837	-0.04004
price	-0.004649	1.000000	0.308350	0.525138	0.702035	0.397293	0.66743
bedrooms	-0.016964	0.308350	1.000000	0.515884	0.576671	0.079532	0.35696
bathrooms	-0.034481	0.525138	0.515884	1.000000	0.754665	0.187737	0.66498
sqft_living	-0.034570	0.702035	0.576671	0.754665	1.000000	0.284611	0.76270
view	-0.001837	0.397293	0.079532	0.187737	0.284611	1.000000	0.25132
grade	-0.040040	0.667434	0.356967	0.664983	0.762704	0.251321	1.00000
sqft_above	-0.027890	0.605567	0.477600	0.685342	0.876597	0.167649	0.75592
sqft_basement	-0.019554	0.323816	0.303093	0.283770	0.435043	0.276947	0.16839
lat	-0.032851	0.307003	-0.008931	0.024573	0.052529	0.006157	0.11408
sqft_living15	-0.031653	0.585379	0.391638	0.568634	0.756420	0.280439	0.71320

```
In [331]: print('`date` corr:', data.corr()[ 'price' ][ 'date' ])
```

```
`date` corr: -0.0046490362965513525
```

Целевой признак price слабо коррелирует с date, удалим date

```
In [332]: data = data.drop(columns='date')
```


4) Проведение корреляционного анализа данных. Формирование промежуточных выводов о возможности построения моделей машинного обучения.

```
In [333]: data.corr()
```

```
Out[333]:
```

	price	bedrooms	bathrooms	sqft_living	view	grade	sqft_above
price	1.000000	0.308350	0.525138	0.702035	0.397293	0.667434	0.605567
bedrooms	0.308350	1.000000	0.515884	0.576671	0.079532	0.356967	0.477600
bathrooms	0.525138	0.515884	1.000000	0.754665	0.187737	0.664983	0.685342
sqft_living	0.702035	0.576671	0.754665	1.000000	0.284611	0.762704	0.876597
view	0.397293	0.079532	0.187737	0.284611	1.000000	0.251321	0.167649
grade	0.667434	0.356967	0.664983	0.762704	0.251321	1.000000	0.755923
sqft_above	0.605567	0.477600	0.685342	0.876597	0.167649	0.755923	1.000000
sqft_basement	0.323816	0.303093	0.283770	0.435043	0.276947	0.168392	-0.051943
lat	0.307003	-0.008931	0.024573	0.052529	0.006157	0.114084	-0.000816
sqft_living15	0.585379	0.391638	0.568634	0.756420	0.280439	0.713202	0.731870

Корреляция была проанализирована выше ([тут](#)-Выбор-признаков,-подходящих-для-построения-моделей.-Кодирование-категориальных-признаков-Масштабирование-данных.-Формирование-вспомогательных-признаков,-улучшающих-качество-моделей.)), выброшены признаки с слабой зависимостью с целевым признаком `price`, он имеет сильную зависимость с `sqft_living`, с остальными – среднюю.

5) Выбор метрик для последующей оценки качества моделей.

Перед нами стоит задача регрессии, необходимо выбрать 3 метрики, подходящие для регрессии.

```
In [334]: from sklearn.metrics import mean_absolute_error, median_absolute_error, r2_score
my_metrics = (mean_absolute_error, median_absolute_error, r2_score)

from sklearn.model_selection import cross_validate, KFold, ShuffleSplit
cross_val_names = ('neg_mean_absolute_error', 'neg_median_absolute_error', 'r2')
```

Выбрана метрика `mean_absolute_error`, поскольку она проста и очевидна, так как отражает разницу в среднем промахе модели ($|y - y'|$ для 1 случая, суммируем N случаев, усредняем, деля на N).

Выбрана метрика `median_absolute_error`, поскольку она устойчива к выбросам (медиана из всех случаев $|y - y'|$).

Выбрана метрика `r2_score`, поскольку она представляет собой универсальную меру зависимости одной случайной величины от множества других.

```
In [335]: def print_metrics(y_true, y_pred):
            for m in my_metrics:
                print('Metric {}: {}'.format(m.__name__, m(y_true, y_pred)))

def print_cross_val_scores(model, X_, y_, cv):
    print(cross_validate(model(), X_, y_, scoring=cross_val_names,
cv=cv, return_train_score=True))
```

6) Выбор наиболее подходящих моделей для решения задачи классификации или регрессии.

Перед нами задача регрессии, выбраны следующие модели:

- LinearRegression
- LinearSVR
- DecisionTreeRegressor
- Ensemble: BaggingRegressor with DecisionTreeRegressor
- Ensemble: Gradient boosting (XGBRegressor from xgboost library)

```
In [336]: from sklearn.linear_model import LinearRegression
from sklearn.svm import LinearSVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import BaggingRegressor
from xgboost import XGBRegressor
```

7) Формирование обучающей и тестовой выборок на основе исходного набора данных.

```
In [337]: from sklearn.model_selection import train_test_split
```

```
In [338]: data.columns
```

```
Out[338]: Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'view', 'grade',
                'sqft_above', 'sqft_basement', 'lat', 'sqft_living15'],
                dtype='object')
```

```
In [339]: test_size = 0.3
state = 42
X, y = data[data.columns[range(1, data.shape[1])]], data[data.columns[[0]]]
xTrain, xTest, yTrain, yTest = train_test_split(X, y, test_size=test_size, random_state=state)
len(xTrain), len(xTest), len(yTrain), len(yTest)
```

```
Out[339]: (15129, 6484, 15129, 6484)
```

```
In [340]: X.columns
```

```
Out[340]: Index(['bedrooms', 'bathrooms', 'sqft_living', 'view', 'grade', 'sqft_above',
                'sqft_basement', 'lat', 'sqft_living15'],
                dtype='object')
```

```
In [341]: y.columns
```

```
Out[341]: Index(['price'], dtype='object')
```

Решение задачи регрессии

8) Построение базового решения (baseline) для выбранных моделей без подбора гиперпараметров.

LinearRegression

```
In [342]: lin_reg = LinearRegression(n_jobs=-1) # no hyperparams  
lin_reg.fit(xTrain, yTrain)
```

```
Out[342]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=-1, normalize=False)
```

```
In [343]: yPredicted_lin_reg = lin_reg.predict(xTest)
```

```
In [344]: print_metrics(yTest, yPredicted_lin_reg)
```

```
Metric mean_absolute_error: 139654.25998331458  
Metric median_absolute_error: 93998.91289327666  
Metric r2_score: 0.6331075681984636
```

LinearSVR

```
In [368]: from sklearn.preprocessing import Normalizer  
normalizerX = Normalizer().fit(X)  
X_n = normalizerX.transform(X)  
  
test_size = 0.3  
state = 42  
xTrain_n, xTest_n = train_test_split(X_n, test_size=test_size, random_state=state)  
print(len(xTrain_n), len(xTest_n))  
  
lin_svr = LinearSVR(C=1.0, max_iter=1000)  
lin_svr.fit(xTrain_n, yTrain.values.ravel())
```

```
15129 6484
```

```
Out[368]: LinearSVR(C=1.0, dual=True, epsilon=0.0, fit_intercept=True,  
intercept_scaling=1.0, loss='epsilon_insensitive', max_iter=1000,  
random_state=None, tol=0.0001, verbose=0)
```

```
In [369]: yPredicted_lin_svr = lin_svr.predict(xTest)
```

```
In [370]: print_metrics(yTest, yPredicted_lin_svr)
```

```
Metric mean_absolute_error: 50270369.62206803  
Metric median_absolute_error: 45918359.11039731  
Metric r2_score: -20096.001594459798
```

DecisionTreeRegressor

```
In [348]: tree = DecisionTreeRegressor(random_state=state) # with default hyperparams
tree.fit(xTrain, yTrain.values.ravel())
```

```
Out[348]: DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                presort=False, random_state=42, splitter='best')
```

```
In [349]: yPredicted_tree = tree.predict(xTest)
```

```
In [350]: print_metrics(yTest, yPredicted_tree)
```

```
Metric mean_absolute_error: 129773.90270923298
Metric median_absolute_error: 70000.5
Metric r2_score: 0.5664820790736267
```

Ensemble: BaggingRegressor with DecisionTreeRegressor

```
In [351]: bagreg_treereg = BaggingRegressor(DecisionTreeRegressor(random_state=state), n_estimators=100)
bagreg_treereg.fit(xTrain, yTrain.values.ravel())
```

```
Out[351]: BaggingRegressor(base_estimator=DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                                                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                                                min_impurity_split=None, min_samples_leaf=1,
                                                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                                                presort=False, random_state=42, splitter='best'),
                             bootstrap=True, bootstrap_features=False, max_features=1.0,
                             max_samples=1.0, n_estimators=100, n_jobs=None, oob_score=False,
                             random_state=None, verbose=0, warm_start=False)
```

```
In [352]: yPredicted_bagreg_treereg = bagreg_treereg.predict(xTest)
```

```
In [353]: print_metrics(yTest, yPredicted_bagreg_treereg)
```

```
Metric mean_absolute_error: 94834.9737203201
Metric median_absolute_error: 51491.19999999995
Metric r2_score: 0.7686480983618686
```

Ensemble: Gradient boosting (XGBRegressor from xgboost library)

```
In [354]: xgbreg_treereg = XGBRegressor(n_jobs=-1)
xgbreg_treereg.fit(xTrain, yTrain.values.ravel())

[14:21:11] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

Out[354]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0,
                        importance_type='gain', learning_rate=0.1, max_delta_step=0,
                        max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
                        n_jobs=-1, nthread=None, objective='reg:linear', random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                        silent=None, subsample=1, verbosity=1)

In [355]: yPredicted_xgbreg_treereg = xgbreg_treereg.predict(xTest)

In [356]: print_metrics(yTest, yPredicted_xgbreg_treereg)

Metric mean_absolute_error: 99196.6039096237
Metric median_absolute_error: 56986.390625
Metric r2_score: 0.7605725228859841
```

9) Подбор гиперпараметров для выбранных моделей.

```
In [359]: from sklearn.model_selection import GridSearchCV, KFold, ShuffleSplit
```

LinearSVR

```
In [381]: parameters = [{'C': np.array(np.arange(0.1, 20.1, 0.1)),
                        'max_iter': np.array([1000, 5000, 10000, 25000, 50000, 100000, 250000, 500000, 1000000])}]
lin_svr_grid = GridSearchCV(LinearSVR(), parameters, cv=ShuffleSplit(n_splits=10, test_size=test_size),
                            scoring='neg_median_absolute_error',
                            n_jobs=-1,
                            verbose=10)
lin_svr_grid.fit(X_n, y.values.ravel())
```

Fitting 10 folds for each of 1800 candidates, totalling 18000 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 2 tasks | elapsed: 0.9s

[Parallel(n_jobs=-1)]: Done	9 tasks	elapsed:	1.0s
[Parallel(n_jobs=-1)]: Done	16 tasks	elapsed:	1.0s
[Parallel(n_jobs=-1)]: Done	25 tasks	elapsed:	1.0s
[Parallel(n_jobs=-1)]: Batch computation too fast (0.1710s.)	Setting batch_size=2.		
[Parallel(n_jobs=-1)]: Done	34 tasks	elapsed:	1.1s
[Parallel(n_jobs=-1)]: Batch computation too fast (0.0703s.)	Setting batch_size=10.		
[Parallel(n_jobs=-1)]: Done	50 tasks	elapsed:	1.1s
[Parallel(n_jobs=-1)]: Done	72 tasks	elapsed:	1.2s
[Parallel(n_jobs=-1)]: Done	202 tasks	elapsed:	1.6s
[Parallel(n_jobs=-1)]: Done	332 tasks	elapsed:	1.9s
[Parallel(n_jobs=-1)]: Done	482 tasks	elapsed:	2.3s
[Parallel(n_jobs=-1)]: Done	632 tasks	elapsed:	2.7s
[Parallel(n_jobs=-1)]: Done	802 tasks	elapsed:	3.1s
[Parallel(n_jobs=-1)]: Done	972 tasks	elapsed:	3.6s
[Parallel(n_jobs=-1)]: Done	1162 tasks	elapsed:	4.2s
[Parallel(n_jobs=-1)]: Done	1352 tasks	elapsed:	4.6s
[Parallel(n_jobs=-1)]: Done	1562 tasks	elapsed:	5.2s
[Parallel(n_jobs=-1)]: Done	1772 tasks	elapsed:	5.8s
[Parallel(n_jobs=-1)]: Done	2002 tasks	elapsed:	6.5s
[Parallel(n_jobs=-1)]: Done	2232 tasks	elapsed:	7.2s
[Parallel(n_jobs=-1)]: Done	2482 tasks	elapsed:	7.8s
[Parallel(n_jobs=-1)]: Done	2732 tasks	elapsed:	8.5s
[Parallel(n_jobs=-1)]: Done	3002 tasks	elapsed:	9.2s
[Parallel(n_jobs=-1)]: Done	3272 tasks	elapsed:	9.9s
[Parallel(n_jobs=-1)]: Done	3562 tasks	elapsed:	10.6s
[Parallel(n_jobs=-1)]: Done	3852 tasks	elapsed:	11.6s
[Parallel(n_jobs=-1)]: Done	4162 tasks	elapsed:	12.9s
[Parallel(n_jobs=-1)]: Done	4472 tasks	elapsed:	13.9s
[Parallel(n_jobs=-1)]: Done	4802 tasks	elapsed:	15.1s
[Parallel(n_jobs=-1)]: Done	5132 tasks	elapsed:	16.3s
[Parallel(n_jobs=-1)]: Done	5482 tasks	elapsed:	17.6s
[Parallel(n_jobs=-1)]: Done	5832 tasks	elapsed:	18.9s
[Parallel(n_jobs=-1)]: Done	6202 tasks	elapsed:	20.3s
[Parallel(n_jobs=-1)]: Done	6572 tasks	elapsed:	21.7s
[Parallel(n_jobs=-1)]: Done	6962 tasks	elapsed:	23.1s
[Parallel(n_jobs=-1)]: Done	7352 tasks	elapsed:	24.5s
[Parallel(n_jobs=-1)]: Done	7762 tasks	elapsed:	26.0s
[Parallel(n_jobs=-1)]: Done	8172 tasks	elapsed:	27.6s
[Parallel(n_jobs=-1)]: Done	8602 tasks	elapsed:	29.2s
[Parallel(n_jobs=-1)]: Done	9032 tasks	elapsed:	30.7s
[Parallel(n_jobs=-1)]: Done	9482 tasks	elapsed:	32.4s
[Parallel(n_jobs=-1)]: Done	9932 tasks	elapsed:	34.1s
[Parallel(n_jobs=-1)]: Done	10402 tasks	elapsed:	36.0s
[Parallel(n_jobs=-1)]: Done	10872 tasks	elapsed:	37.8s
[Parallel(n_jobs=-1)]: Done	11362 tasks	elapsed:	39.6s
[Parallel(n_jobs=-1)]: Done	11852 tasks	elapsed:	41.5s
[Parallel(n_jobs=-1)]: Done	12362 tasks	elapsed:	43.4s
[Parallel(n_jobs=-1)]: Done	12872 tasks	elapsed:	45.3s
[Parallel(n_jobs=-1)]: Done	13402 tasks	elapsed:	47.4s
[Parallel(n_jobs=-1)]: Done	13932 tasks	elapsed:	49.6s
[Parallel(n_jobs=-1)]: Done	14482 tasks	elapsed:	51.7s
[Parallel(n_jobs=-1)]: Done	15032 tasks	elapsed:	53.9s
[Parallel(n_jobs=-1)]: Done	15602 tasks	elapsed:	56.0s
[Parallel(n_jobs=-1)]: Done	16172 tasks	elapsed:	58.2s
[Parallel(n_jobs=-1)]: Done	16762 tasks	elapsed:	1.0min

```
[Parallel(n_jobs=-1)]: Done 17352 tasks      | elapsed:  1.0min  
[Parallel(n_jobs=-1)]: Done 18000 out of 18000 | elapsed:  1.1min  
finished
```

```
Out[381]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=None, test_
size=0.3, train_size=None),
                error_score='raise-deprecating',
                estimator=LinearSVR(C=1.0, dual=True, epsilon=0.0, fit_intercept=True,
                intercept_scaling=1.0, loss='epsilon_insensitive', max_iter=1000,
                random_state=None, tol=0.0001, verbose=0),
                fit_params=None, iid='warn', n_jobs=-1,
                param_grid=[{'C': array([ 0.1,  0.2, ..., 19.9, 20. ]), 'max_iter': array([ 1000,  5000, 10000, 25000, 50000, 100000, 250000, 500000, 1000000])}],
                pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                scoring='neg_median_absolute_error', verbose=10)
```

```
In [390]: lin_svr_grid.best_estimator_
```

```
Out[390]: LinearSVR(C=20.000000000000004, dual=True, epsilon=0.0, fit_intercept=True,
                intercept_scaling=1.0, loss='epsilon_insensitive', max_iter=100000,
                random_state=None, tol=0.0001, verbose=0)
```

```
In [391]: lin_svr_grid.best_score_
```

```
Out[391]: -146027.5578859402
```

```
In [392]: lin_svr_grid.best_params_
```

```
Out[392]: {'C': 20.000000000000004, 'max_iter': 100000}
```

DecisionTreeRegressor

```
In [396]: parameters = [{'random_state': np.array([state]),
                        'max_depth': np.array([None, 10, 50, 100, 500, 1000,
                        5000, 10000]),
                        'min_samples_split': np.array(range(2, 11)),
                        'min_samples_leaf': np.array(range(1, 11))
                        }]
tree_grid = GridSearchCV(DecisionTreeRegressor(), parameters, cv=ShuffleSplit(n_splits=10, test_size=0.2),
                        scoring='neg_median_absolute_error',
                        n_jobs=-1,
                        verbose=10)
tree_grid.fit(X, y.values.ravel())
```

Fitting 10 folds for each of 720 candidates, totalling 7200 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 2 tasks	elapsed: 0.7s
[Parallel(n_jobs=-1)]: Done 9 tasks	elapsed: 0.9s
[Parallel(n_jobs=-1)]: Done 16 tasks	elapsed: 0.9s
[Parallel(n_jobs=-1)]: Done 25 tasks	elapsed: 1.2s
[Parallel(n_jobs=-1)]: Done 34 tasks	elapsed: 1.3s
[Parallel(n_jobs=-1)]: Done 45 tasks	elapsed: 1.5s
[Parallel(n_jobs=-1)]: Done 56 tasks	elapsed: 1.6s
[Parallel(n_jobs=-1)]: Done 69 tasks	elapsed: 1.9s
[Parallel(n_jobs=-1)]: Done 82 tasks	elapsed: 2.2s
[Parallel(n_jobs=-1)]: Done 97 tasks	elapsed: 2.5s
[Parallel(n_jobs=-1)]: Done 112 tasks	elapsed: 2.8s
[Parallel(n_jobs=-1)]: Done 129 tasks	elapsed: 3.1s
[Parallel(n_jobs=-1)]: Done 146 tasks	elapsed: 3.4s
[Parallel(n_jobs=-1)]: Done 165 tasks	elapsed: 3.8s
[Parallel(n_jobs=-1)]: Done 184 tasks	elapsed: 4.0s
[Parallel(n_jobs=-1)]: Done 205 tasks	elapsed: 4.4s
[Parallel(n_jobs=-1)]: Done 226 tasks	elapsed: 4.8s
[Parallel(n_jobs=-1)]: Done 249 tasks	elapsed: 5.2s
[Parallel(n_jobs=-1)]: Done 272 tasks	elapsed: 5.5s
[Parallel(n_jobs=-1)]: Done 297 tasks	elapsed: 5.9s
[Parallel(n_jobs=-1)]: Done 322 tasks	elapsed: 6.3s
[Parallel(n_jobs=-1)]: Done 349 tasks	elapsed: 6.7s
[Parallel(n_jobs=-1)]: Done 376 tasks	elapsed: 7.1s
[Parallel(n_jobs=-1)]: Done 405 tasks	elapsed: 7.6s
[Parallel(n_jobs=-1)]: Done 434 tasks	elapsed: 8.0s
[Parallel(n_jobs=-1)]: Done 465 tasks	elapsed: 8.5s
[Parallel(n_jobs=-1)]: Done 496 tasks	elapsed: 8.9s
[Parallel(n_jobs=-1)]: Done 529 tasks	elapsed: 9.4s
[Parallel(n_jobs=-1)]: Done 562 tasks	elapsed: 9.9s
[Parallel(n_jobs=-1)]: Done 597 tasks	elapsed: 10.4s
[Parallel(n_jobs=-1)]: Done 632 tasks	elapsed: 10.9s
[Parallel(n_jobs=-1)]: Done 669 tasks	elapsed: 11.4s
[Parallel(n_jobs=-1)]: Done 706 tasks	elapsed: 12.0s
[Parallel(n_jobs=-1)]: Done 745 tasks	elapsed: 12.6s
[Parallel(n_jobs=-1)]: Done 784 tasks	elapsed: 13.1s
[Parallel(n_jobs=-1)]: Done 825 tasks	elapsed: 13.7s
[Parallel(n_jobs=-1)]: Done 866 tasks	elapsed: 14.3s
[Parallel(n_jobs=-1)]: Done 909 tasks	elapsed: 14.8s
[Parallel(n_jobs=-1)]: Done 952 tasks	elapsed: 15.4s
[Parallel(n_jobs=-1)]: Batch computation too fast (0.1975s.) Setting batch_size=2.	
[Parallel(n_jobs=-1)]: Done 997 tasks	elapsed: 16.0s
[Parallel(n_jobs=-1)]: Done 1085 tasks	elapsed: 17.0s
[Parallel(n_jobs=-1)]: Done 1179 tasks	elapsed: 18.0s
[Parallel(n_jobs=-1)]: Done 1273 tasks	elapsed: 19.1s
[Parallel(n_jobs=-1)]: Done 1371 tasks	elapsed: 20.3s
[Parallel(n_jobs=-1)]: Done 1469 tasks	elapsed: 21.5s
[Parallel(n_jobs=-1)]: Done 1571 tasks	elapsed: 22.7s
[Parallel(n_jobs=-1)]: Done 1673 tasks	elapsed: 23.9s
[Parallel(n_jobs=-1)]: Done 1779 tasks	elapsed: 25.1s
[Parallel(n_jobs=-1)]: Done 1885 tasks	elapsed: 27.2s
[Parallel(n_jobs=-1)]: Done 1995 tasks	elapsed: 29.3s
[Parallel(n_jobs=-1)]: Done 2105 tasks	elapsed: 31.3s
[Parallel(n_jobs=-1)]: Done 2219 tasks	elapsed: 33.1s
[Parallel(n_jobs=-1)]: Done 2333 tasks	elapsed: 35.0s

```

[Parallel(n_jobs=-1)]: Done 2451 tasks      | elapsed: 36.7s
[Parallel(n_jobs=-1)]: Done 2569 tasks      | elapsed: 38.6s
[Parallel(n_jobs=-1)]: Done 2691 tasks      | elapsed: 40.2s
[Parallel(n_jobs=-1)]: Done 2813 tasks      | elapsed: 42.6s
[Parallel(n_jobs=-1)]: Done 2939 tasks      | elapsed: 44.8s
[Parallel(n_jobs=-1)]: Done 3065 tasks      | elapsed: 46.9s
[Parallel(n_jobs=-1)]: Done 3195 tasks      | elapsed: 49.1s
[Parallel(n_jobs=-1)]: Done 3325 tasks      | elapsed: 51.2s
[Parallel(n_jobs=-1)]: Done 3459 tasks      | elapsed: 53.3s
[Parallel(n_jobs=-1)]: Done 3593 tasks      | elapsed: 55.6s
[Parallel(n_jobs=-1)]: Done 3731 tasks      | elapsed: 58.5s
[Parallel(n_jobs=-1)]: Done 3869 tasks      | elapsed: 1.0min
[Parallel(n_jobs=-1)]: Done 4011 tasks      | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 4153 tasks      | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 4299 tasks      | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 4445 tasks      | elapsed: 1.2min
[Parallel(n_jobs=-1)]: Done 4595 tasks      | elapsed: 1.2min
[Parallel(n_jobs=-1)]: Done 4745 tasks      | elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 4899 tasks      | elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 5053 tasks      | elapsed: 1.4min
[Parallel(n_jobs=-1)]: Done 5211 tasks      | elapsed: 1.4min
[Parallel(n_jobs=-1)]: Done 5369 tasks      | elapsed: 1.4min
[Parallel(n_jobs=-1)]: Done 5531 tasks      | elapsed: 1.5min
[Parallel(n_jobs=-1)]: Done 5693 tasks      | elapsed: 1.5min
[Parallel(n_jobs=-1)]: Done 5859 tasks      | elapsed: 1.6min
[Parallel(n_jobs=-1)]: Done 6025 tasks      | elapsed: 1.6min
[Parallel(n_jobs=-1)]: Done 6195 tasks      | elapsed: 1.7min
[Parallel(n_jobs=-1)]: Done 6365 tasks      | elapsed: 1.7min
[Parallel(n_jobs=-1)]: Done 6539 tasks      | elapsed: 1.8min
[Parallel(n_jobs=-1)]: Done 6713 tasks      | elapsed: 1.8min
[Parallel(n_jobs=-1)]: Done 6891 tasks      | elapsed: 1.9min
[Parallel(n_jobs=-1)]: Done 7069 tasks      | elapsed: 1.9min
[Parallel(n_jobs=-1)]: Done 7185 out of 7200 | elapsed: 2.0min re
maining: 0.2s
[Parallel(n_jobs=-1)]: Done 7200 out of 7200 | elapsed: 2.0min fi
nished

```

```

Out[396]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=None, test_
size=0.2, train_size=None),
                    error_score='raise-deprecating',
                    estimator=DecisionTreeRegressor(criterion='mse', max_depth=
None, max_features=None,
                    max_leaf_nodes=None, min_impurity_decrease=0.0,
                    min_impurity_split=None, min_samples_leaf=1,
                    min_samples_split=2, min_weight_fraction_leaf=0.0,
                    presort=False, random_state=None, splitter='best'),
                    fit_params=None, iid='warn', n_jobs=-1,
                    param_grid=[{'random_state': array([42]), 'max_depth': arra
y([None, 10, 50, 100, 500, 1000, 5000, 10000], dtype=object), 'min
_samples_split': array([ 2, 3, 4, 5, 6, 7, 8, 9, 10]), 'min
_samples_leaf': array([ 1, 2, 3, 4, 5, 6, 7, 8, 9, 10])}],
                    pre_dispatch='2*n_jobs', refit=True, return_train_score='wa
rn',
                    scoring='neg_median_absolute_error', verbose=10)

```


Fitting 10 folds for each of 20 candidates, totalling 200 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done	2 tasks	elapsed:	1.5s
[Parallel(n_jobs=-1)]: Done	9 tasks	elapsed:	1.7s
[Parallel(n_jobs=-1)]: Done	16 tasks	elapsed:	2.2s
[Parallel(n_jobs=-1)]: Done	25 tasks	elapsed:	3.3s
[Parallel(n_jobs=-1)]: Done	34 tasks	elapsed:	4.8s
[Parallel(n_jobs=-1)]: Done	45 tasks	elapsed:	7.8s
[Parallel(n_jobs=-1)]: Done	56 tasks	elapsed:	10.8s
[Parallel(n_jobs=-1)]: Done	69 tasks	elapsed:	17.5s
[Parallel(n_jobs=-1)]: Done	82 tasks	elapsed:	23.3s
[Parallel(n_jobs=-1)]: Done	97 tasks	elapsed:	31.6s
[Parallel(n_jobs=-1)]: Done	112 tasks	elapsed:	41.5s
[Parallel(n_jobs=-1)]: Done	129 tasks	elapsed:	55.2s
[Parallel(n_jobs=-1)]: Done	146 tasks	elapsed:	1.2min
[Parallel(n_jobs=-1)]: Done	165 tasks	elapsed:	1.5min
[Parallel(n_jobs=-1)]: Done	184 tasks	elapsed:	1.8min
[Parallel(n_jobs=-1)]: Done	200 out of 200	elapsed:	2.1min finished

```
Out[401]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=None, test_size=0.2, train_size=None),
                        error_score='raise-deprecating',
```

```
estimator=BaggingRegressor(base_estimator=DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                                                                    max_leaf_nodes=None, min_impurity_decrease=0.0,
                                                                    min_impurity_split=None, min_samples_leaf=1,
                                                                    min_samples_split=2, min_weight_fraction_leaf=0.0,
                                                                    ...stimators=10, n_jobs=None, oob_score=False,
                                                                    random_state=None, verbose=0, warm_start=False),
fit_params=None, iid='warn', n_jobs=-1,
param_grid=[{'n_estimators': array([ 1,  6, 11, 16, 21, 26,
31, 36, 41, 46, 51, 56, 61, 66, 71, 76, 81,
86, 91, 96])}],
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring='neg_median_absolute_error', verbose=10)
```

```
In [403]: bagreg_grid.best_estimator_
```

```
Out[403]: BaggingRegressor(base_estimator=DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                                                                    max_leaf_nodes=None, min_impurity_decrease=0.0,
                                                                    min_impurity_split=None, min_samples_leaf=1,
                                                                    min_samples_split=2, min_weight_fraction_leaf=0.0,
                                                                    presort=False, random_state=42, splitter='best'),
bootstrap=True, bootstrap_features=False, max_features=1.0,
max_samples=1.0, n_estimators=71, n_jobs=None, oob_score=False,
random_state=None, verbose=0, warm_start=False)
```

```
In [404]: bagreg_grid.best_score_
```

```
Out[404]: -50558.99342723005
```

```
In [405]: bagreg_grid.best_params_
```

```
Out[405]: {'n_estimators': 71}
```

Ensemble: Gradient boosting (XGBRegressor from xgboost library)

```
In [406]: parameters = [{"colsample_bytree": [1.0], "min_child_weight": [0.8,
1.0, 1.2],
                        'max_depth': range(3, 11), 'n_estimators': [25, 50,
75, 100]}]
xgbreg_grid = GridSearchCV(XGBRegressor(), parameters, cv=ShuffleSp
lit(n_splits=10, test_size=0.2),
                           scoring='neg_median_absolute_error',
                           n_jobs=-1,
                           verbose=10,
                           )
xgbreg_grid.fit(X, y.values.ravel())
```

Fitting 10 folds for each of 96 candidates, totalling 960 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done	2 tasks	elapsed:	0.4s
[Parallel(n_jobs=-1)]: Done	9 tasks	elapsed:	0.8s
[Parallel(n_jobs=-1)]: Done	16 tasks	elapsed:	1.3s
[Parallel(n_jobs=-1)]: Done	25 tasks	elapsed:	2.8s
[Parallel(n_jobs=-1)]: Done	34 tasks	elapsed:	4.3s
[Parallel(n_jobs=-1)]: Done	45 tasks	elapsed:	5.4s
[Parallel(n_jobs=-1)]: Done	56 tasks	elapsed:	6.2s
[Parallel(n_jobs=-1)]: Done	69 tasks	elapsed:	8.3s
[Parallel(n_jobs=-1)]: Done	82 tasks	elapsed:	9.7s
[Parallel(n_jobs=-1)]: Done	97 tasks	elapsed:	11.2s
[Parallel(n_jobs=-1)]: Done	112 tasks	elapsed:	13.6s
[Parallel(n_jobs=-1)]: Done	129 tasks	elapsed:	15.4s
[Parallel(n_jobs=-1)]: Done	146 tasks	elapsed:	17.7s
[Parallel(n_jobs=-1)]: Done	165 tasks	elapsed:	21.1s
[Parallel(n_jobs=-1)]: Done	184 tasks	elapsed:	23.5s
[Parallel(n_jobs=-1)]: Done	205 tasks	elapsed:	27.0s
[Parallel(n_jobs=-1)]: Done	226 tasks	elapsed:	29.6s
[Parallel(n_jobs=-1)]: Done	249 tasks	elapsed:	33.4s
[Parallel(n_jobs=-1)]: Done	272 tasks	elapsed:	39.0s
[Parallel(n_jobs=-1)]: Done	297 tasks	elapsed:	42.8s
[Parallel(n_jobs=-1)]: Done	322 tasks	elapsed:	48.3s
[Parallel(n_jobs=-1)]: Done	349 tasks	elapsed:	53.7s
[Parallel(n_jobs=-1)]: Done	376 tasks	elapsed:	58.3s
[Parallel(n_jobs=-1)]: Done	405 tasks	elapsed:	1.1min
[Parallel(n_jobs=-1)]: Done	434 tasks	elapsed:	1.2min
[Parallel(n_jobs=-1)]: Done	465 tasks	elapsed:	1.3min
[Parallel(n_jobs=-1)]: Done	496 tasks	elapsed:	1.4min
[Parallel(n_jobs=-1)]: Done	529 tasks	elapsed:	1.6min
[Parallel(n_jobs=-1)]: Done	562 tasks	elapsed:	1.7min
[Parallel(n_jobs=-1)]: Done	597 tasks	elapsed:	1.9min
[Parallel(n_jobs=-1)]: Done	632 tasks	elapsed:	2.0min
[Parallel(n_jobs=-1)]: Done	669 tasks	elapsed:	2.2min
[Parallel(n_jobs=-1)]: Done	706 tasks	elapsed:	2.4min
[Parallel(n_jobs=-1)]: Done	745 tasks	elapsed:	2.6min
[Parallel(n_jobs=-1)]: Done	784 tasks	elapsed:	2.8min
[Parallel(n_jobs=-1)]: Done	825 tasks	elapsed:	3.0min
[Parallel(n_jobs=-1)]: Done	866 tasks	elapsed:	3.2min
[Parallel(n_jobs=-1)]: Done	909 tasks	elapsed:	3.5min
[Parallel(n_jobs=-1)]: Done	960 out of 960	elapsed:	3.8min finished

[14:58:22] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

```
Out[406]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=None, test_
size=0.2, train_size=None),
                      error_score='raise-deprecating',
                      estimator=XGBRegressor(base_score=0.5, booster='gbtree', co
lsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, gamma=0,
                      importance_type='gain', learning_rate=0.1, max_delta_step=0
,
                      max_depth=3, min_child_weight=1, missing=None, n_estimators
=100,
                      n_jobs=1, nthread=None, objective='reg:linear', random_stat
e=0,
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                      silent=None, subsample=1, verbosity=1),
                      fit_params=None, iid='warn', n_jobs=-1,
                      param_grid=[{'colsample_bytree': [1.0], 'min_child_weight':
[0.8, 1.0, 1.2], 'max_depth': range(3, 11), 'n_estimators': [25, 5
0, 75, 100]}],
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='wa
rn',
                      scoring='neg_median_absolute_error', verbose=10)
```

```
In [407]: xgbreg_grid.best_estimator_
```

```
Out[407]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1
,
                      colsample_bynode=1, colsample_bytree=1.0, gamma=0,
                      importance_type='gain', learning_rate=0.1, max_delta_step=0
,
                      max_depth=10, min_child_weight=0.8, missing=None, n_estimat
ors=50,
                      n_jobs=1, nthread=None, objective='reg:linear', random_stat
e=0,
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                      silent=None, subsample=1, verbosity=1)
```

```
In [408]: xgbreg_grid.best_score_
```

```
Out[408]: -49993.25625
```

```
In [409]: xgbreg_grid.best_params_
```

```
Out[409]: {'colsample_bytree': 1.0,
           'max_depth': 10,
           'min_child_weight': 0.8,
           'n_estimators': 50}
```

10) Повторение пункта 8 для найденных оптимальных значений гиперпараметров. Сравнение качества полученных моделей с качеством baseline-моделей.

LinearSVR

```
In [410]: lin_svr_grid.best_estimator_.fit(xTrain_n, yTrain.values.ravel())

yPredicted_lin_svr_new = lin_svr_grid.best_estimator_.predict(xTest_n)
print('old:')
print_metrics(yTest, yPredicted_lin_svr)

print('\nnew:')
print_metrics(yTest, yPredicted_lin_svr_new)

old:
Metric mean_absolute_error: 50270369.62206803
Metric median_absolute_error: 45918359.11039731
Metric r2_score: -20096.001594459798

new:
Metric mean_absolute_error: 262271.02348911285
Metric median_absolute_error: 147450.00928251055
Metric r2_score: -0.3718708004825009
```

DecisionTreeRegressor

```
In [414]: tree_grid.best_estimator_.fit(xTrain, yTrain.values.ravel())

yPredicted_tree_new = tree_grid.best_estimator_.predict(xTest)
print('old:')
print_metrics(yTest, yPredicted_tree)

print('\nnew:')
print_metrics(yTest, yPredicted_tree_new)

old:
Metric mean_absolute_error: 129773.90270923298
Metric median_absolute_error: 70000.5
Metric r2_score: 0.5664820790736267

new:
Metric mean_absolute_error: 109969.3788295769
Metric median_absolute_error: 59306.32438316394
Metric r2_score: 0.6733312258690927
```

Ensemble: BaggingRegressor with DecisionTreeRegressor


```
In [412]: bagreg_grid.best_estimator_.fit(xTrain, yTrain.values.ravel())

yPredicted_bagreg_treereg_new = bagreg_grid.best_estimator_.predict(
    xTest)
print('old:')
print_metrics(yTest, yPredicted_bagreg_treereg)

print('\nnew:')
print_metrics(yTest, yPredicted_bagreg_treereg_new)

old:
Metric mean_absolute_error: 94834.9737203201
Metric median_absolute_error: 51491.19999999995
Metric r2_score: 0.7686480983618686

new:
Metric mean_absolute_error: 94800.05277012683
Metric median_absolute_error: 51570.542253521126
Metric r2_score: 0.7697305739843688
```

Ensemble: Gradient boosting (XGBRegressor from xgboost library)

```
In [413]: xgbreg_grid.best_estimator_.fit(xTrain, yTrain.values.ravel())

yPredicted_xgbreg_treereg_new = xgbreg_grid.best_estimator_.predict(
    xTest)
print('old:')
print_metrics(yTest, yPredicted_xgbreg_treereg)

print('\nnew:')
print_metrics(yTest, yPredicted_xgbreg_treereg_new)

[14:59:08] WARNING: src/objective/regression_obj.cu:152: reg:linear
is now deprecated in favor of reg:squarederror.
old:
Metric mean_absolute_error: 99196.6039096237
Metric median_absolute_error: 56986.390625
Metric r2_score: 0.7605725228859841

new:
Metric mean_absolute_error: 93738.3927745219
Metric median_absolute_error: 51429.953125
Metric r2_score: 0.7570938895100083
```

Графическая реализация

11) Формирование выводов о качестве построенных моделей на основе выбранных метрик.

```
In [415]: from sklearn.model_selection import learning_curve, validation_curve
```

```
In [416]: def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                                n_jobs=-1, train_sizes=np.linspace(.1, 1.0,
                                5)):
    """
    Generate a simple plot of the test and training learning curve.

    Parameters
    -----
    estimator : object type that implements the "fit" and "predict"
    methods
        An object of that type which is cloned for each validation.

    title : string
        Title for the chart.

    X : array-like, shape (n_samples, n_features)
        Training vector, where n_samples is the number of samples a
    nd
        n_features is the number of features.

    y : array-like, shape (n_samples) or (n_samples, n_features), o
    ptional
        Target relative to X for classification or regression;
        None for unsupervised learning.

    ylim : tuple, shape (ymin, ymax), optional
        Defines minimum and maximum yvalues plotted.

    cv : int, cross-validation generator or an iterable, optional
        Determines the cross-validation splitting strategy.
        Possible inputs for cv are:
        - None, to use the default 3-fold cross-validation,
        - integer, to specify the number of folds.
        - :term:`CV splitter`,
        - An iterable yielding (train, test) splits as arrays of
    indices.

        For integer/None inputs, if ``y`` is binary or multiclass,
        :class:`StratifiedKFold` used. If the estimator is not a cl
    assifier
        or if ``y`` is neither binary nor multiclass, :class:`KFold`
        is used.

        Refer :ref:`User Guide <cross_validation>` for the various
        cross-validators that can be used here.
```

```

    n_jobs : int or None, optional (default=None)
        Number of jobs to run in parallel.
        ``None`` means 1 unless in a :obj:`joblib.parallel_backend`
context.
        ``-1`` means using all processors. See :term:`Glossary <n_j
obs>`
        for more details.

    train_sizes : array-like, shape (n_ticks,), dtype float or int
        Relative or absolute numbers of training examples that will
be used to
        generate the learning curve. If the dtype is float, it is r
egarded as a
        fraction of the maximum size of the training set (that is d
etermined
        by the selected validation method), i.e. it has to be withi
n (0, 1].
        Otherwise it is interpreted as absolute sizes of the traini
ng sets.
        Note that for classification the number of samples usually
have to
        be big enough to contain at least one sample from each clas
s.
        (default: np.linspace(0.1, 1.0, 5))
    """
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_si
zes)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    plt.grid()

    plt.fill_between(train_sizes, train_scores_mean - train_scores_
std,
                    train_scores_mean + train_scores_std, alpha=0.
1,
                    color="r")
    plt.fill_between(train_sizes, test_scores_mean - test_scores_st
d,
                    test_scores_mean + test_scores_std, alpha=0.1,
color="g")
    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Cross-validation score")

    plt.legend(loc="best")
    return plt

```

```

In [417]: def plot_validation_curve(estimator, title, X, y,
                                     param_name, param_range, cv,
                                     scoring="accuracy"):

    train_scores, test_scores = validation_curve(
        estimator, X, y, param_name=param_name, param_range=param_r
ange,
        cv=cv, scoring=scoring,
        n_jobs=-1
    )
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)

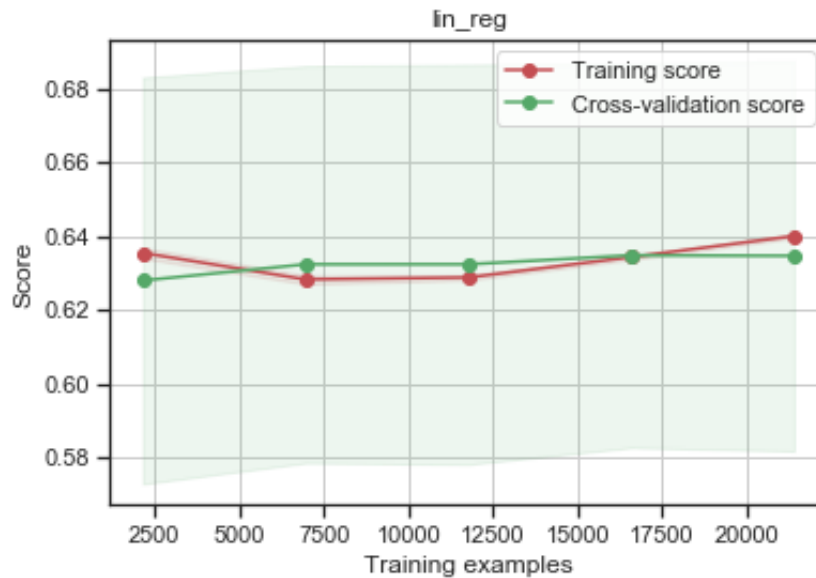
    plt.title(title)
    plt.xlabel(param_name)
    plt.ylabel("Score")
    plt.ylim(0.0, 1.1)
    lw = 2
    plt.plot(param_range, train_scores_mean, label="Training score"
,
            color="darkorange", lw=lw)
    plt.fill_between(param_range, train_scores_mean - train_scores_
std,
                    train_scores_mean + train_scores_std, alpha=0.
2,
                    color="darkorange", lw=lw)
    plt.plot(param_range, test_scores_mean, label="Cross-validation
score",
            color="navy", lw=lw)
    plt.fill_between(param_range, test_scores_mean - test_scores_st
d,
                    test_scores_mean + test_scores_std, alpha=0.2,
                    color="navy", lw=lw)
    plt.legend(loc="best")
    return plt

```

LinearRegression

```
In [433]: plot_learning_curve(lin_reg, 'lin_reg',  
                               X, y.values.ravel(), cv=KFold(n_splits=100))
```

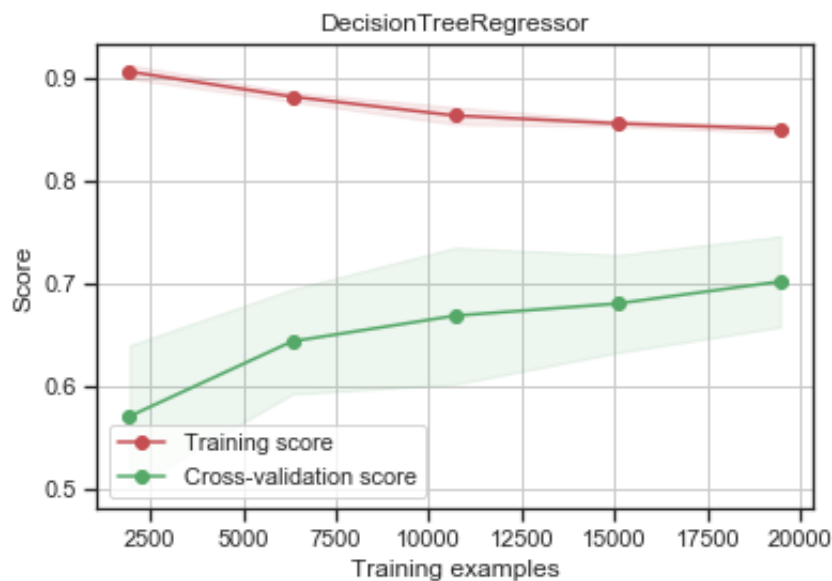
```
Out[433]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.ven  
v/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```



DecisionTreeRegressor


```
In [430]: plot_learning_curve(tree_grid.best_estimator_, 'DecisionTreeRegressor',  
                               X, y.values.ravel(), cv=KFold(n_splits=10))
```

```
Out[430]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.ven  
v/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```




```
In [453]: plt.plot(range(720), tree_grid.cv_results_['mean_train_score'])
```

```
Out[453]: [<matplotlib.lines.Line2D at 0x17099bcc0>]
```



```
In [454]: plt.plot(range(720), tree_grid.cv_results_['mean_test_score'])
```


```
Out[454]: [<matplotlib.lines.Line2D at 0x1706b1160>]
```



```
In [440]: parameters_tree = [{'random_state': np.array([state]),
                              'max_depth': np.array([None, 10, 50, 100, 500,
1000, 5000, 10000])},
                              'min_samples_split': np.array(range(2, 11)),
                              'min_samples_leaf': np.array(range(1, 11))
                              }]


plot_validation_curve(DecisionTreeRegressor(random_state=state), 'D
ecisionTreeRegressor',
                      X, y.values.ravel(),
                      param_name='max_depth', param_range=[10, 50,
100, 500, 1000, 5000, 10000],
                      cv=KFold(n_splits=10), scoring="neg_median_ab
solute_error")
```

```
Out[440]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.ven
v/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```



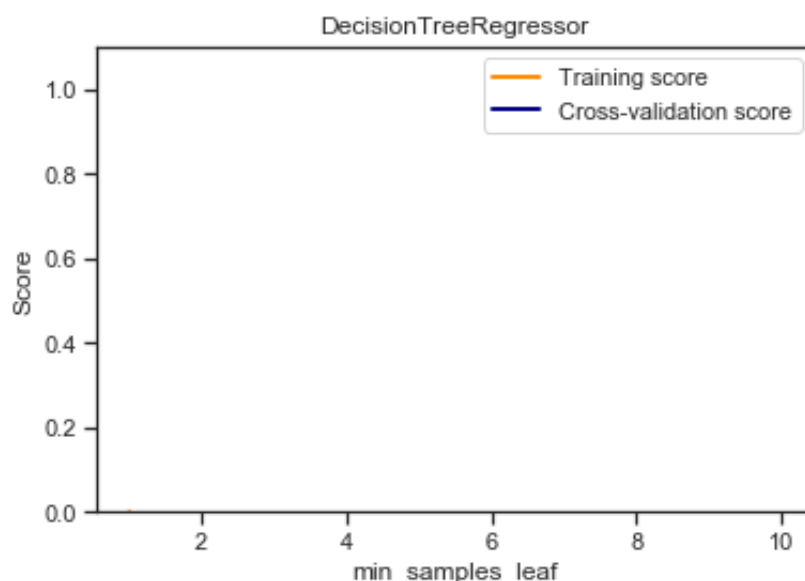
```
In [436]: plot_validation_curve(DecisionTreeRegressor(random_state=state), 'DecisionTreeRegressor',
                                X, y.values.ravel(),
                                param_name='min_samples_split', param_range=range(2, 11),
                                cv=KFold(n_splits=10), scoring="neg_median_absolute_error")
```

```
Out[436]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.venv/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```



```
In [438]: plot_validation_curve(DecisionTreeRegressor(random_state=state), 'DecisionTreeRegressor',
                                X, y.values.ravel(),
                                param_name='min_samples_leaf', param_range=range(1, 11),
                                cv=KFold(n_splits=10), scoring="neg_median_absolute_error")
```

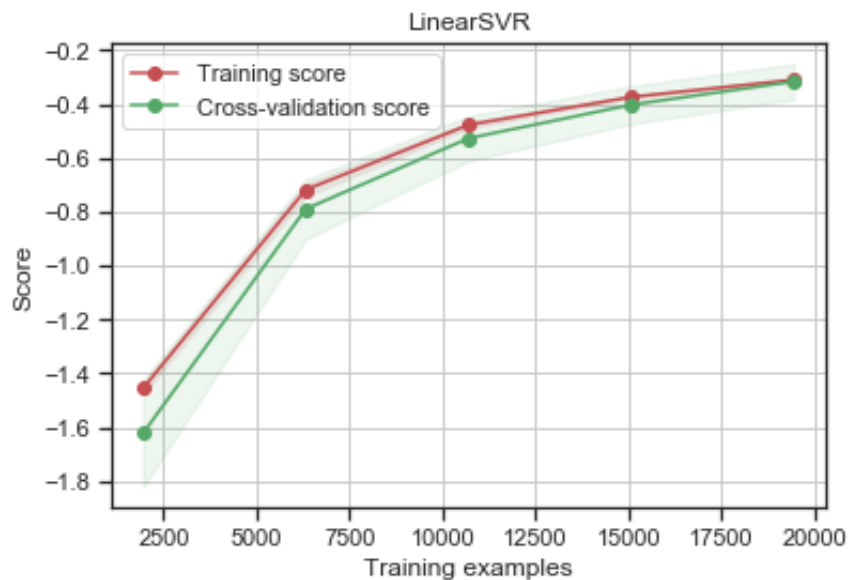
```
Out[438]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.venv/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```



LinearSVR


```
In [442]: plot_learning_curve(lin_svr_grid.best_estimator_, 'LinearSVR',  
                             X_n, y.values.ravel(), cv=KFold(n_splits=10))
```

```
Out[442]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.ven  
v/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```




```
In [459]: plt.plot(range(1800), lin_svr_grid.cv_results_['mean_train_score'])
```

```
Out[459]: [<matplotlib.lines.Line2D at 0x170abe7f0>]
```



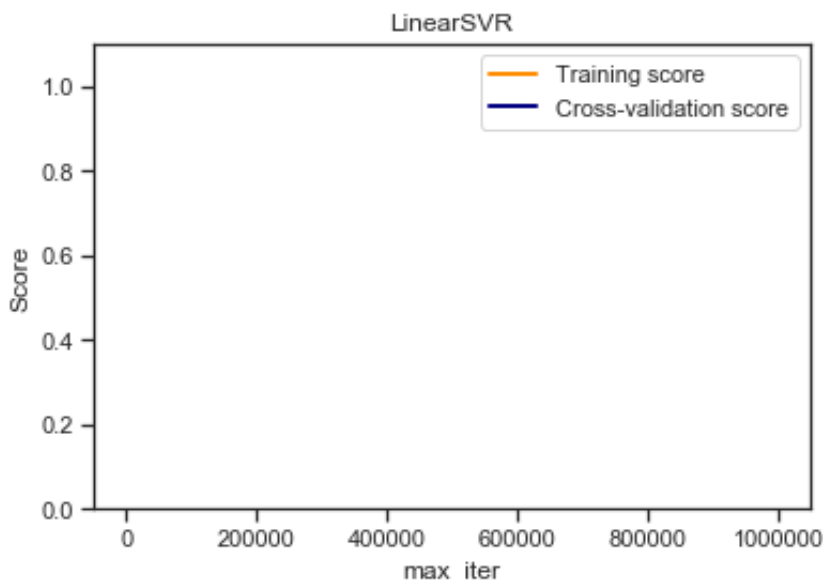
```
In [460]: plt.plot(range(1800), lin_svr_grid.cv_results_['mean_test_score'])
```

```
Out[460]: [<matplotlib.lines.Line2D at 0x16f538080>]
```



```
In [443]: parameters_lin_svr = {'max_iter': np.array([1000, 5000, 10000, 25000, 50000, 100000, 250000, 500000, 1000000])}
plot_validation_curve(LinearSVR(), 'LinearSVR',
                        x, y.values.ravel(),
                        param_name='max_iter', param_range=[1000, 5000, 10000, 25000, 50000, 100000, 250000, 500000, 1000000],
                        cv=KFold(n_splits=10), scoring="neg_median_absolute_error")
```

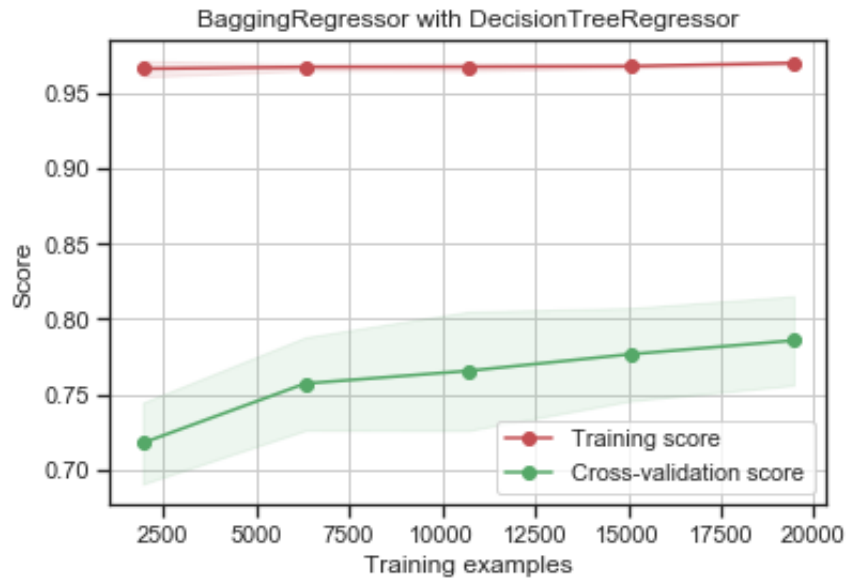
```
Out[443]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.venv/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```



Ensemble: BaggingRegressor with DecisionTreeRegressor


```
In [445]: plot_learning_curve(bagreg_grid.best_estimator_, 'BaggingRegressor  
with DecisionTreeRegressor',  
                             X, y.values.ravel(), cv=KFold(n_splits=10))
```

```
Out[445]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.ven  
v/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```




```
In [462]: plt.plot(range(20), bagreg_grid.cv_results_['mean_train_score'])
```

```
Out[462]: [<matplotlib.lines.Line2D at 0x170ccd240>]
```



```
In [463]: plt.plot(range(20), bagreg_grid.cv_results_['mean_test_score'])
```

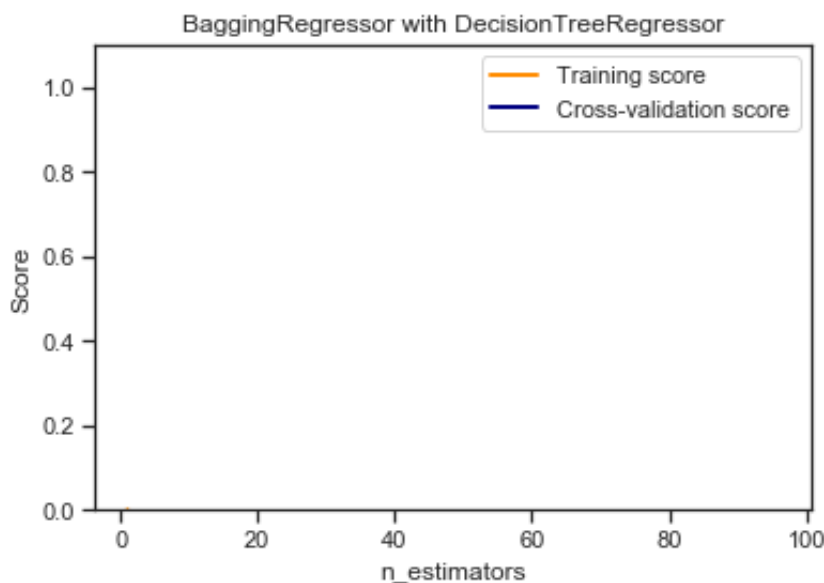
```
Out[463]: [<matplotlib.lines.Line2D at 0x170204400>]
```



```
In [469]: parameters_bagging = [{'n_estimators': np.array(range(1, 101, 5))}]

plot_validation_curve(BaggingRegressor(DecisionTreeRegressor(random
_state=state)), 'BaggingRegressor with DecisionTreeRegressor',
                    X, y.values.ravel(),
                    param_name='n_estimators', param_range=np.arr
ay(range(1, 101, 5)),
                    cv=KFold(n_splits=10), scoring="neg_median_ab
solute_error")
```

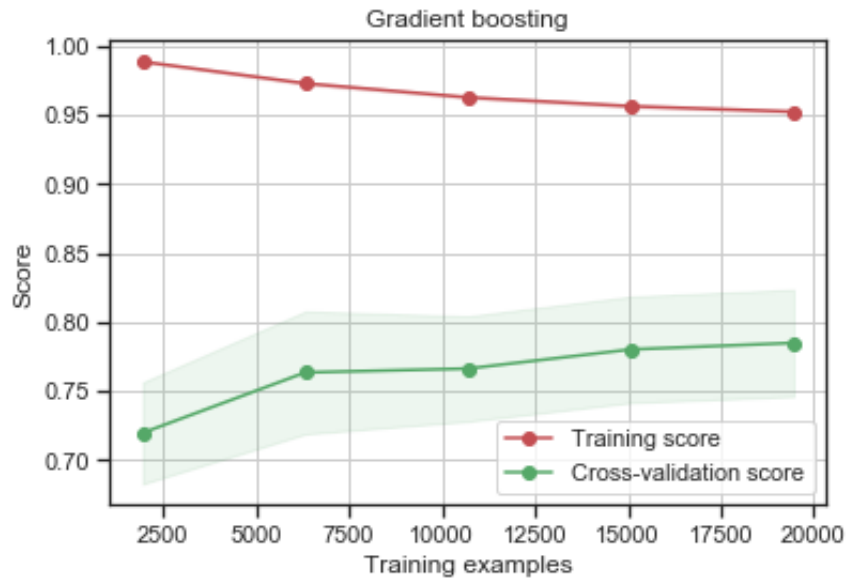
```
Out[469]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.ven
v/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```



Ensemble: Gradient boosting (XGBRegressor from xgboost library)


```
In [449]: plot_learning_curve(xgbreg_grid.best_estimator_, 'Gradient boosting',
                              X, y.values.ravel(), cv=KFold(n_splits=10))
```

```
Out[449]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.venv/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```




```
In [467]: plt.plot(range(96), xgbreg_grid.cv_results_['mean_train_score'])
```

```
Out[467]: [<matplotlib.lines.Line2D at 0x171193748>]
```



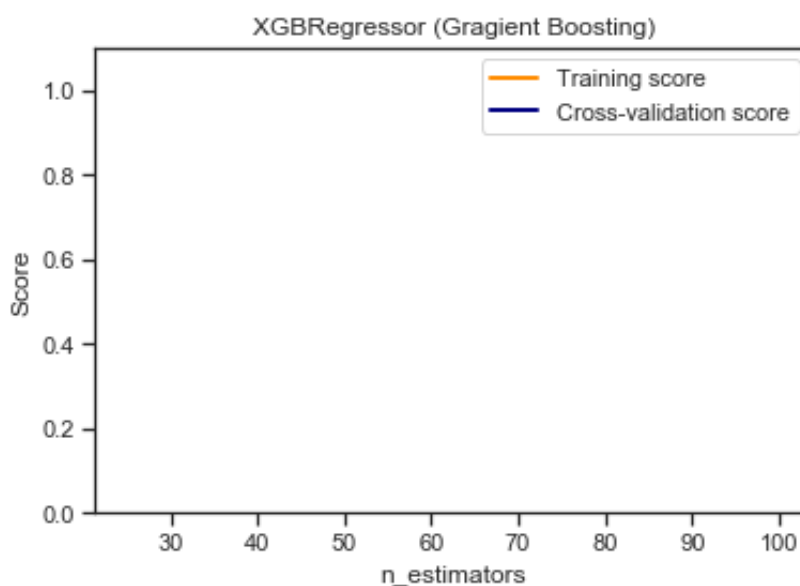
```
In [468]: plt.plot(range(96), xgbreg_grid.cv_results_['mean_test_score'])
```

```
Out[468]: [<matplotlib.lines.Line2D at 0x16015d080>]
```



```
In [450]: parameters_xgb = [{"colsample_bytree": [1.0], "min_child_weight": [0.8, 1.0, 1.2],  
                             'max_depth': range(3, 11), 'n_estimators': [25, 50,  
                             75, 100]}]  
  
plot_validation_curve(XGBRegressor(), 'XGBRegressor (Gragient Boosting)',  
                      x, y.values.ravel(),  
                      param_name='n_estimators', param_range=[25, 50, 75, 100],  
                      cv=KFold(n_splits=10), scoring="neg_median_absolute_error")
```

```
Out[450]: <module 'matplotlib.pyplot' from '/Users/artiom.andreev/Study/.venv/lib/python3.7/site-packages/matplotlib/pyplot.py'>
```



```
In [ ]:
```

Заключение

Таким образом, внедрение технологий машинного обучения может помочь решению задач, которые возникают во всех сферах жизни человека, в предсказании стоимости различных объектов в зависимости от их свойств, времени, географического положения итд. На этом датасете самые хорошие результаты по всем метрикам качества показали методы ансамбль градиентного бустинга на базе XGBoost, Бэггинг на базе дерева решений для регрессии и само дерево решений для регрессии.

Список литературы

1. Лекции Гапанюка Ю.Е. [Электронный ресурс]. – Электрон. дан. - URL: https://github.com/ugapanyuk/ml_course/wiki/COURSE_TMO (дата обращения: 10.05.2019)
2. Breast Cancer Wisconsin (Diagnostic) Data Set [Электронный ресурс]. – Электрон. дан. - URL: <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data> (дата обращения: 10.05.2019)
3. Открытый курс машинного обучения. Тема 4. Линейные модели классификации и регрессии [Электронный ресурс]. – Электрон. дан. - URL: <https://habr.com/ru/company/ods/blog/323890/> (дата обращения: 10.05.2019)
4. Руководство для начинающих [Электронный ресурс]. – Электрон. дан. - URL: <https://mlbootcamp.ru/article/tutorial/> (дата обращения: 10.05.2019)