

**Московский государственный технический
университет им. Н.Э. Баумана.**

Факультет «Информатика и управление»

Кафедра ИУ5. Курс «Технологии машинного обучения»

Отчет по лабораторной работе №6:
«Ансамбли моделей машинного обучения»

Выполнил:
студент группы ИУ5-62
Андреев Артем

Проверил:

Подпись и дата:

Подпись и дата:

Москва, 2019 г.

```
In [1]: import numpy as np
import pandas as pd
pd.set_option('display.max.rows', 1000)
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style='ticks')
```

Подготовка датасета

```
In [2]: # House Sales in King County, USA
# Predict house price using regression
data = pd.read_csv('data/kc_house_data.csv')
data.shape
```

Out[2]: (21613, 21)

```
In [3]: data.head()
```

Out[3]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basem
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0	0	...	7	1180	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0	0	...	7	2170	
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0	0	...	6	770	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0	0	...	7	1050	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0	0	...	8	1680	

5 rows × 21 columns

```
In [4]: data = data.drop(columns=['id', 'date', ])
data.corr()
```

Out[4]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_bas
price	1.000000	0.308350	0.525138	0.702035	0.089661	0.256794	0.266369	0.397293	0.036362	0.667434	0.605567	0.323816
bedrooms	0.308350	1.000000	0.515884	0.576671	0.031703	0.175429	-0.006582	0.079532	0.028472	0.356967	0.477600	0.303093
bathrooms	0.525138	0.515884	1.000000	0.754665	0.087740	0.500653	0.063744	0.187737	-0.124982	0.664983	0.685342	0.283770
sqft_living	0.702035	0.576671	0.754665	1.000000	0.172826	0.353949	0.103818	0.284611	-0.058753	0.762704	0.876597	0.435043
sqft_lot	0.089661	0.031703	0.087740	0.172826	1.000000	-0.005201	0.021604	0.074710	-0.008958	0.113621	0.183512	0.015286
floors	0.256794	0.175429	0.500653	0.353949	-0.005201	1.000000	0.023698	0.029444	-0.263768	0.458183	0.523885	-0.245705
waterfront	0.266369	-0.006582	0.063744	0.103818	0.021604	0.023698	1.000000	0.401857	0.016653	0.082775	0.072075	0.080588
view	0.397293	0.079532	0.187737	0.284611	0.074710	0.029444	0.401857	1.000000	0.045990	0.251321	0.167649	0.276947
condition	0.036362	0.028472	-0.124982	-0.058753	-0.008958	-0.263768	0.016653	0.045990	1.000000	-0.144674	-0.158214	0.174105
grade	0.667434	0.356967	0.664983	0.762704	0.113621	0.458183	0.082775	0.251321	-0.144674	1.000000	0.755923	0.168392
sqft_above	0.605567	0.477600	0.685342	0.876597	0.183512	0.523885	0.072075	0.167649	-0.158214	0.755923	1.000000	-0.051943
sqft_basement	0.323816	0.303093	0.283770	0.435043	0.015286	-0.245705	0.080588	0.276947	0.174105	0.168392	-0.051943	1.000000
yr_built	0.054012	0.154178	0.506019	0.318049	0.053080	0.489319	-0.026161	-0.053440	-0.361417	0.446963	0.423898	-0.026161
yr_renovated	0.126434	0.018841	0.050739	0.055363	0.007644	0.006338	0.092885	0.103917	-0.060618	0.014414	0.023285	0.006338
zipcode	-0.053203	-0.152668	-0.203866	-0.199430	-0.129574	-0.059121	0.030285	0.084827	0.003026	-0.184862	-0.261190	0.030285
lat	0.307003	-0.008931	0.024573	0.052529	-0.085683	0.049614	-0.014274	0.006157	-0.014941	0.114084	-0.000816	0.006157
long	0.021626	0.129473	0.223042	0.240223	0.229521	0.125419	-0.041910	-0.078400	-0.106500	0.198372	0.343803	-0.041910
sqft_living15	0.585379	0.391638	0.568634	0.756420	0.144608	0.279885	0.086463	0.280439	-0.092824	0.713202	0.731870	0.279885
sqft_lot15	0.082447	0.029244	0.087175	0.183286	0.718557	-0.011269	0.030703	0.072575	-0.003406	0.119248	0.194050	0.030703

```
In [5]: data.isnull().sum()
```

```
Out[5]: price           0  
bedrooms              0  
bathrooms            0  
sqft_living          0  
sqft_lot             0  
floors               0  
waterfront           0  
view                 0  
condition            0  
grade                0  
sqft_above           0  
sqft_basement        0  
yr_built             0  
yr_renovated         0  
zipcode              0  
lat                  0  
long                 0  
sqft_living15        0  
sqft_lot15           0  
dtype: int64
```

```
In [6]: data.dtypes
```

```
Out[6]: price           float64  
bedrooms              int64  
bathrooms            float64  
sqft_living           int64  
sqft_lot              int64  
floors                float64  
waterfront            int64  
view                  int64  
condition             int64  
grade                 int64  
sqft_above            int64  
sqft_basement         int64  
yr_built              int64  
yr_renovated          int64  
zipcode               int64  
lat                   float64  
long                  float64  
sqft_living15         int64  
sqft_lot15            int64  
dtype: object
```

```
In [7]: # пропусков нет, разделим на обучающую и тестовую выборку  
from sklearn.model_selection import train_test_split
```

```
In [8]: # перед этим разделим исходный датасет на 2: один содержит независимые параметры, другой — зависимый (price)  
X, y = data[data.columns[range(1, 19)]], data[data.columns[[0]]]
```

```
In [9]: X.dtypes
```

```
Out[9]: bedrooms          int64  
bathrooms          float64  
sqft_living         int64  
sqft_lot            int64  
floors              float64  
waterfront          int64  
view                int64  
condition           int64  
grade              int64  
sqft_above          int64  
sqft_basement       int64  
yr_built            int64  
yr_renovated        int64  
zipcode             int64  
lat                 float64  
long                float64  
sqft_living15       int64  
sqft_lot15          int64  
dtype: object
```

```
In [10]: y.dtypes
```

```
Out[10]: price          float64  
dtype: object
```

```
In [11]: test_size = 0.2  
state = 42  
xTrain, xTest, yTrain, yTest = train_test_split(X, y, test_size=test_size, random_state=state)  
len(xTrain), len(xTest), len(yTrain), len(yTest)
```

```
Out[11]: (17290, 4323, 17290, 4323)
```

Обучение моделей

BaggingRegressor with DecisionTreeRegressor

```
In [12]: from sklearn.ensemble import BaggingRegressor
         from sklearn.tree import DecisionTreeRegressor
```

```
state = 42
```

```
bagreg_treereg = BaggingRegressor(DecisionTreeRegressor(random_state=state), n_estimators=100)
bagreg_treereg.fit(X, y.values.ravel())
```

```
Out[12]: BaggingRegressor(base_estimator=DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
presort=False, random_state=42, splitter='best'),
bootstrap=True, bootstrap_features=False, max_features=1.0,
max_samples=1.0, n_estimators=100, n_jobs=None, oob_score=False,
random_state=None, verbose=0, warm_start=False)
```

```
In [13]: bagreg_treeregTrain = bagreg_treereg.predict(xTrain)
```

```
In [14]: bagreg_treeregTest = bagreg_treereg.predict(xTest)
```

```
In [64]: # оценим качество модели регрессии
         from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

         # 1) mean absolute error – средняя абсолютная ошибка
         print('mean_absolute_error (train): {}'.format(mean_absolute_error(yTrain, bagreg_treeregTrain)))
         print('mean_absolute_error (test): {}'.format(mean_absolute_error(yTest, bagreg_treeregTest)))
```

```
mean_absolute_error (train): 25312.712859152005
mean_absolute_error (test): 26819.885470475383
```

```
In [31]: from xgboost import XGBRegressor
```

```
xgbreg_treereg = XGBRegressor(n_jobs=-1)
```

```
In [32]: xgbreg_treereg.fit(X, y.values.ravel())
```

```
[23:06:53] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squareerror.
```

```
Out[32]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, gamma=0,
                      importance_type='gain', learning_rate=0.1, max_delta_step=0,
                      max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
                      n_jobs=-1, nthread=None, objective='reg:linear', random_state=0,
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                      silent=None, subsample=1, verbosity=1)
```

```
In [33]: xgbreg_treeregTrain = xgbreg_treereg.predict(xTrain)
```

```
In [34]: xgbreg_treeregTest = xgbreg_treereg.predict(xTest)
```

```
In [65]: # оценим качество модели регрессии
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# 1) mean absolute error – средняя абсолютная ошибка
print('mean_absolute_error (train): {}'.format(mean_absolute_error(yTrain, xgbreg_treeregTrain)))
print('mean_absolute_error (test): {}'.format(mean_absolute_error(yTest, xgbreg_treeregTest)))

mean_absolute_error (train): 72965.081767821
mean_absolute_error (test): 75977.6707147814
```

подбор одного гиперпараметра с использованием GridSearchCV и кросс-валидации

[illegible]

Fitting 10 folds for each of 100 candidates, totalling 1000 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.  
[Parallel(n_jobs=-1)]: Done   2 tasks      | elapsed:   1.3s  
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:   1.5s  
[Parallel(n_jobs=-1)]: Done  16 tasks      | elapsed:   1.7s  
[Parallel(n_jobs=-1)]: Done  25 tasks      | elapsed:   2.4s  
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:   3.2s  
[Parallel(n_jobs=-1)]: Done  45 tasks      | elapsed:   4.5s  
[Parallel(n_jobs=-1)]: Done  56 tasks      | elapsed:   5.8s  
[Parallel(n_jobs=-1)]: Done  69 tasks      | elapsed:   8.2s  
[Parallel(n_jobs=-1)]: Done  82 tasks      | elapsed:  10.7s  
[Parallel(n_jobs=-1)]: Done  97 tasks      | elapsed:  14.3s  
[Parallel(n_jobs=-1)]: Done 112 tasks      | elapsed:  18.2s  
[Parallel(n_jobs=-1)]: Done 129 tasks      | elapsed:  23.5s  
[Parallel(n_jobs=-1)]: Done 146 tasks      | elapsed:  29.1s  
[Parallel(n_jobs=-1)]: Done 165 tasks      | elapsed:  36.8s  
[Parallel(n_jobs=-1)]: Done 184 tasks      | elapsed:  44.6s  
[Parallel(n_jobs=-1)]: Done 205 tasks      | elapsed:  55.2s  
[Parallel(n_jobs=-1)]: Done 226 tasks      | elapsed:  1.1min  
[Parallel(n_jobs=-1)]: Done 249 tasks      | elapsed:  1.3min  
[Parallel(n_jobs=-1)]: Done 272 tasks      | elapsed:  1.6min  
[Parallel(n_jobs=-1)]: Done 297 tasks      | elapsed:  1.9min  
[Parallel(n_jobs=-1)]: Done 322 tasks      | elapsed:  2.2min  
[Parallel(n_jobs=-1)]: Done 349 tasks      | elapsed:  2.6min  
[Parallel(n_jobs=-1)]: Done 376 tasks      | elapsed:  2.9min  
[Parallel(n_jobs=-1)]: Done 405 tasks      | elapsed:  3.4min  
[Parallel(n_jobs=-1)]: Done 434 tasks      | elapsed:  3.9min  
[Parallel(n_jobs=-1)]: Done 465 tasks      | elapsed:  4.4min  
[Parallel(n_jobs=-1)]: Done 496 tasks      | elapsed:  5.0min  
[Parallel(n_jobs=-1)]: Done 529 tasks      | elapsed:  5.7min  
[Parallel(n_jobs=-1)]: Done 562 tasks      | elapsed:  6.4min  
[Parallel(n_jobs=-1)]: Done 597 tasks      | elapsed:  7.2min  
[Parallel(n_jobs=-1)]: Done 632 tasks      | elapsed:  8.0min  
[Parallel(n_jobs=-1)]: Done 669 tasks      | elapsed:  9.0min  
[Parallel(n_jobs=-1)]: Done 706 tasks      | elapsed:  9.9min  
[Parallel(n_jobs=-1)]: Done 745 tasks      | elapsed: 11.1min  
[Parallel(n_jobs=-1)]: Done 784 tasks      | elapsed: 12.2min  
[Parallel(n_jobs=-1)]: Done 825 tasks      | elapsed: 13.6min  
[Parallel(n_jobs=-1)]: Done 866 tasks      | elapsed: 14.9min  
[Parallel(n_jobs=-1)]: Done 909 tasks      | elapsed: 16.5min  
[Parallel(n_jobs=-1)]: Done 952 tasks      | elapsed: 18.1min  
[Parallel(n_jobs=-1)]: Done 1000 out of 1000 | elapsed: 19.9min finished
```

```
Out[41]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=None, test_size=0.2, train_size=None),
    error_score='raise-deprecating',
    estimator=BaggingRegressor(base_estimator=DecisionTreeRegressor(criterion='mse', max_depth=None, max_
features=None,
    max_leaf_nodes=None, min_impurity_decrease=0.0,
    min_impurity_split=None, min_samples_leaf=1,
    min_samples_split=2, min_weight_fraction_leaf=0.0,
    ...stimators=10, n_jobs=None, oob_score=False,
    random_state=None, verbose=0, warm_start=False),
    fit_params=None, iid='warn', n_jobs=-1,
    param_grid=[{'n_estimators': array([ 1,  2, ..., 99, 100])}],
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='neg_mean_absolute_error', verbose=10)
```

```
In [42]: bagreg_grid.best_estimator_
```

```
Out[42]: BaggingRegressor(base_estimator=DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
    max_leaf_nodes=None, min_impurity_decrease=0.0,
    min_impurity_split=None, min_samples_leaf=1,
    min_samples_split=2, min_weight_fraction_leaf=0.0,
    presort=False, random_state=42, splitter='best'),
    bootstrap=True, bootstrap_features=False, max_features=1.0,
    max_samples=1.0, n_estimators=88, n_jobs=None, oob_score=False,
    random_state=None, verbose=0, warm_start=False)
```

```
In [44]: bagreg_grid.best_score_
```

```
Out[44]: -69083.26590519112
```

```
In [45]: bagreg_grid.best_params_
```

```
Out[45]: {'n_estimators': 88}
```

```
In [49]: parameters = [{"colsample_bytree": [1.0], "min_child_weight": [0.8, 1.0, 1.2],  
                        'max_depth': range(3, 11), 'n_estimators': [25, 50, 75, 100]}]  
  
xgbreg_grid = GridSearchCV(XGBRegressor(), parameters, cv=ShuffleSplit(n_splits=10, test_size=0.2), scoring=  
    'neg_mean_absolute_error',  
                             n_jobs=-1,  
                             verbose=10,  
                             )  
  
xgbreg_grid.fit(X, y.values.ravel())
```

Fitting 10 folds for each of 96 candidates, totalling 960 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.  
[Parallel(n_jobs=-1)]: Done   2 tasks      | elapsed:   1.8s  
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:   2.5s  
[Parallel(n_jobs=-1)]: Done  16 tasks      | elapsed:   3.2s  
[Parallel(n_jobs=-1)]: Done  25 tasks      | elapsed:   5.8s  
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:   8.4s  
[Parallel(n_jobs=-1)]: Done  45 tasks      | elapsed:  10.4s  
[Parallel(n_jobs=-1)]: Done  56 tasks      | elapsed:  11.8s  
[Parallel(n_jobs=-1)]: Done  69 tasks      | elapsed:  15.6s  
[Parallel(n_jobs=-1)]: Done  82 tasks      | elapsed:  17.9s  
[Parallel(n_jobs=-1)]: Done  97 tasks      | elapsed:  20.4s  
[Parallel(n_jobs=-1)]: Done 112 tasks      | elapsed:  24.7s  
[Parallel(n_jobs=-1)]: Done 129 tasks      | elapsed:  27.5s  
[Parallel(n_jobs=-1)]: Done 146 tasks      | elapsed:  31.8s  
[Parallel(n_jobs=-1)]: Done 165 tasks      | elapsed:  37.9s  
[Parallel(n_jobs=-1)]: Done 184 tasks      | elapsed:  42.4s  
[Parallel(n_jobs=-1)]: Done 205 tasks      | elapsed:  48.8s  
[Parallel(n_jobs=-1)]: Done 226 tasks      | elapsed:  53.7s  
[Parallel(n_jobs=-1)]: Done 249 tasks      | elapsed:  1.0min  
[Parallel(n_jobs=-1)]: Done 272 tasks      | elapsed:  1.2min  
[Parallel(n_jobs=-1)]: Done 297 tasks      | elapsed:  1.3min  
[Parallel(n_jobs=-1)]: Done 322 tasks      | elapsed:  1.5min  
[Parallel(n_jobs=-1)]: Done 349 tasks      | elapsed:  1.6min  
[Parallel(n_jobs=-1)]: Done 376 tasks      | elapsed:  1.7min  
[Parallel(n_jobs=-1)]: Done 405 tasks      | elapsed:  2.0min  
[Parallel(n_jobs=-1)]: Done 434 tasks      | elapsed:  2.2min  
[Parallel(n_jobs=-1)]: Done 465 tasks      | elapsed:  2.3min  
[Parallel(n_jobs=-1)]: Done 496 tasks      | elapsed:  2.6min  
[Parallel(n_jobs=-1)]: Done 529 tasks      | elapsed:  2.8min  
[Parallel(n_jobs=-1)]: Done 562 tasks      | elapsed:  3.1min  
[Parallel(n_jobs=-1)]: Done 597 tasks      | elapsed:  3.3min  
[Parallel(n_jobs=-1)]: Done 632 tasks      | elapsed:  3.6min  
[Parallel(n_jobs=-1)]: Done 669 tasks      | elapsed:  3.9min  
[Parallel(n_jobs=-1)]: Done 706 tasks      | elapsed:  4.2min  
[Parallel(n_jobs=-1)]: Done 745 tasks      | elapsed:  4.6min  
[Parallel(n_jobs=-1)]: Done 784 tasks      | elapsed:  4.9min  
[Parallel(n_jobs=-1)]: Done 825 tasks      | elapsed:  5.3min  
[Parallel(n_jobs=-1)]: Done 866 tasks      | elapsed:  5.7min  
[Parallel(n_jobs=-1)]: Done 909 tasks      | elapsed:  6.3min  
[Parallel(n_jobs=-1)]: Done 960 out of 960 | elapsed:  6.8min finished
```

[00:03:01] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

```
Out[49]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=None, test_size=0.2, train_size=None),
    error_score='raise-deprecating',
    estimator=XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0,
    importance_type='gain', learning_rate=0.1, max_delta_step=0,
    max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
    n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, subsample=1, verbosity=1),
    fit_params=None, iid='warn', n_jobs=-1,
    param_grid=[{'colsample_bytree': [1.0], 'min_child_weight': [0.8, 1.0, 1.2], 'max_depth': range(3, 1
    1), 'n_estimators': [25, 50, 75, 100]}],
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='neg_mean_absolute_error', verbose=10)
```

```
In [50]: xgbreg_grid.best_estimator_
```

```
Out[50]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bynode=1, colsample_bytree=1.0, gamma=0,
    importance_type='gain', learning_rate=0.1, max_delta_step=0,
    max_depth=9, min_child_weight=0.8, missing=None, n_estimators=100,
    n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, subsample=1, verbosity=1)
```

```
In [51]: xgbreg_grid.best_score_
```

```
Out[51]: -65791.59962085068
```

```
In [52]: xgbreg_grid.best_params_
```

```
Out[52]: {'colsample_bytree': 1.0,
    'max_depth': 9,
    'min_child_weight': 0.8,
    'n_estimators': 100}
```


Обучение с новыми параметрами

```
In [62]: bagreg_grid.best_estimator_.fit(xTrain, yTrain.values.ravel())

bagreg_treeregTrainNew = bagreg_grid.best_estimator_.predict(xTrain)
bagreg_treeregTestNew = bagreg_grid.best_estimator_.predict(xTest)

print('mean_absolute_error (train): {}'.format(mean_absolute_error(yTrain, bagreg_treeregTrain)))
print('mean_absolute_error (test): {}'.format(mean_absolute_error(yTest, bagreg_treeregTest)))

print('New mean_absolute_error (train): {}'.format(mean_absolute_error(yTrain, bagreg_treeregTrainNew)))
print('New mean_absolute_error (test): {}'.format(mean_absolute_error(yTest, bagreg_treeregTestNew)))

mean_absolute_error (train): 25312.712859152005
mean_absolute_error (test): 26819.885470475383
New mean_absolute_error (train): 26035.191542553624
New mean_absolute_error (test): 73352.7241413403
```

```
In [63]: xgbreg_grid.best_estimator_.fit(xTrain, yTrain.values.ravel())

xgbreg_treeregTrainNew = xgbreg_grid.best_estimator_.predict(xTrain)
xgbreg_treeregTestNew = xgbreg_grid.best_estimator_.predict(xTest)

print('mean_absolute_error (train): {}'.format(mean_absolute_error(yTrain, xgbreg_treeregTrain)))
print('mean_absolute_error (test): {}'.format(mean_absolute_error(yTest, xgbreg_treeregTest)))

print('New mean_absolute_error (train): {}'.format(mean_absolute_error(yTrain, xgbreg_treeregTrainNew)))
print('New mean_absolute_error (test): {}'.format(mean_absolute_error(yTest, xgbreg_treeregTestNew)))

[00:14:27] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:square
derror.
mean_absolute_error (train): 72965.081767821
mean_absolute_error (test): 75977.6707147814
New mean_absolute_error (train): 33665.23249439705
New mean_absolute_error (test): 69095.68306876012
```

Улучшение для xgbreg_grid

In []: